

News Article Classification NLP Project

Goal of this project is to build a classification model that can automatically categorize news articles into different predefined categories.

In today's digital landscape, news articles are generated and shared at an unprecedented scale. For platforms like news aggregators, media outlets, and social networks, automatically classifying articles into categories such as **sports**, **politics**, and **technology** is essential for improving content organization and recommendation systems.

This project aims to build a robust machine learning model that can accurately classify news articles into predefined categories based on their textual content. By automating this process, we enable faster and smarter content delivery, helping readers access relevant information aligned with their interests.

Key objectives:

- Cleaning and preparing the text data.
- Preprocess and vectorize text data using TF-IDF and Word Embedding.
- Train and evaluate classification models (Logistic Regression, Naive Bayes, SVM).
- Deploy a prediction pipeline that returns the most likely category for any given article.

Data Collection and Preprocessing

Importing Libraries

```
In [3]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

```
In [4]: #importing library for NLP project

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer

import string
import re
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
nltk.download('punkt')
nltk.download('punkt_tab')
```

```
nltk.download('wordnet')

import warnings
warnings.filterwarnings('ignore')

[nltk_data] Downloading package stopwords to
[nltk_data]      C:\Users\vanda\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]      C:\Users\vanda\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to
[nltk_data]      C:\Users\vanda\AppData\Roaming\nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data]      C:\Users\vanda\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

In [5]: # Reading Dataset
df=pd.read_excel("data_news.xlsx")

In [6]: datanews_df=df.copy()

In [7]: datanews_df.head()

	category	headline	links	short_descrip
0	WELLNESS	143 Miles in 35 Days: Lessons Learned	https://www.huffingtonpost.com/entry/running-l...	Resting is part of training confirmed
1	WELLNESS	Talking to Yourself: Crazy or Helpful?	https://www.huffingtonpost.com/entry/talking-t...	Think of talking to yourself as a tool to c
2	WELLNESS	Crenezumab: Trial Will Gauge Whether Alzheimer...	https://www.huffingtonpost.com/entry/crenezuma...	The clinical trial is taking place in the United States
3	WELLNESS	Oh, What a Difference She Made	https://www.huffingtonpost.com/entry/meaningfu...	If you want to be busy, keep time to be productive
4	WELLNESS	Green Superfoods	https://www.huffingtonpost.com/entry/green-sup...	First, the news: Soda ban cornered because

In [8]: datanews_df.shape # displaying shape of Datanews dataset

Out[8]: (50000, 5)

In [9]: datanews_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   category        50000 non-null   object  
 1   headline         50000 non-null   object  
 2   links            50000 non-null   object  
 3   short_description 49994 non-null   object  
 4   keywords          47294 non-null   object  
dtypes: object(5)
memory usage: 1.9+ MB
```

In [10]: `# Finding Duplicates and removing them`
`print(datanews_df.duplicated().sum())`
`datanews_df.drop_duplicates(inplace=True)`

4251

In [11]: `print(datanews_df.duplicated().sum())`

0

Removed Duplicates

Handles missing values in short_description and keywords columns

In [12]: `datanews_df.shape`

Out[12]: (45749, 5)

In [13]: `# Handling missing value`
`datanews_df['keywords']=datanews_df['keywords'].fillna('unknown')`
`datanews_df['short_description']=datanews_df['short_description'].fillna('No Des')`

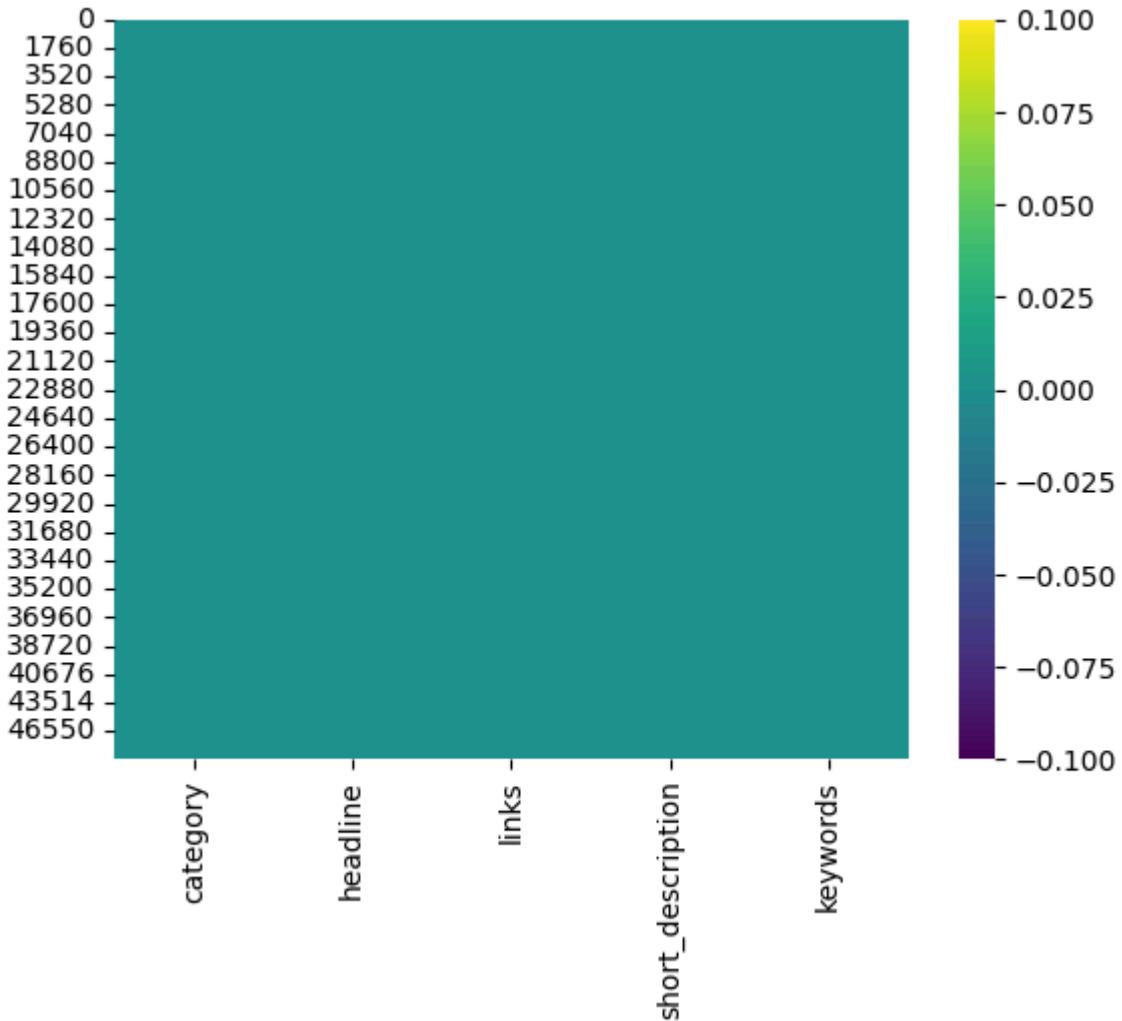
In [14]: `datanews_df.isnull().sum()`

Out[14]:

category	0
headline	0
links	0
short_description	0
keywords	0

dtype: int64

In [15]: `sns.heatmap(datanews_df.isnull(), cmap="viridis")`
`plt.show()`

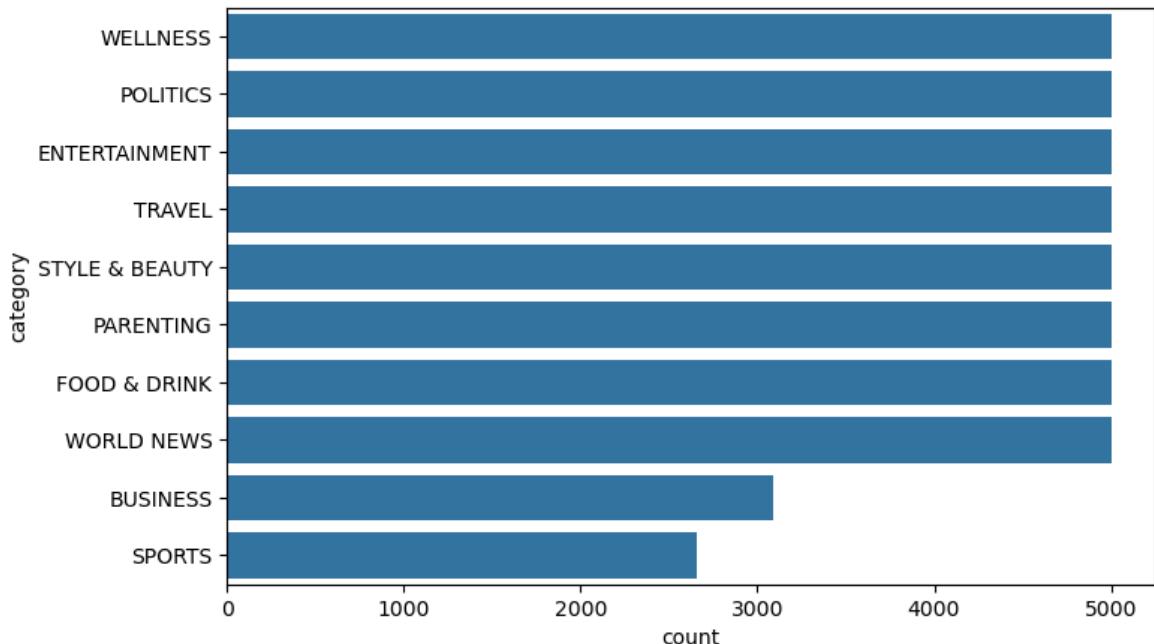


No missing data present in dataset

```
In [16]: datanews_df['category'].value_counts()
```

```
Out[16]: category
WELLNESS      5000
POLITICS      5000
ENTERTAINMENT  5000
TRAVEL         5000
STYLE & BEAUTY 5000
PARENTING      5000
FOOD & DRINK   5000
WORLD NEWS     5000
BUSINESS        3091
SPORTS          2658
Name: count, dtype: int64
```

```
In [17]: plt.figure(figsize=(8,5))
sns.countplot(datanews_df['category'])
plt.show()
```



Dataset is imbalanced.

Most categories (like WELLNESS, POLITICS, ENTERTAINMENT) have close to 5000 articles.

BUSINESS and SPORTS have fewer — around 3091 and 2658, respectively.

In [18]: `datanews_df.head()`

	category	headline	links	short_descrip
0	WELLNESS	143 Miles in 35 Days: Lessons Learned	https://www.huffingtonpost.com/entry/running-l...	Resting is part of training confirmed
1	WELLNESS	Talking to Yourself: Crazy or Helpful?	https://www.huffingtonpost.com/entry/talking-t...	Think of talking to yourself as a way to c
2	WELLNESS	Crenezumab: Trial Will Gauge Whether Alzheimer...	https://www.huffingtonpost.com/entry/crenezuma...	The clinical trial for United States
3	WELLNESS	Oh, What a Difference She Made	https://www.huffingtonpost.com/entry/meaningfu...	If you want to be busy, keep trying to be p
4	WELLNESS	Green Superfoods	https://www.huffingtonpost.com/entry/green-sup...	First, the news: Soda ban could be p

Analysing Length of Headline Short_description and keywords

```
In [19]: # Combine headline, description, and keywords
datanews_df['combine_text'] = (
    datanews_df['headline'].astype(str) + " " +
    datanews_df['short_description'].astype(str) + " " +
    datanews_df['keywords'].astype(str)
)

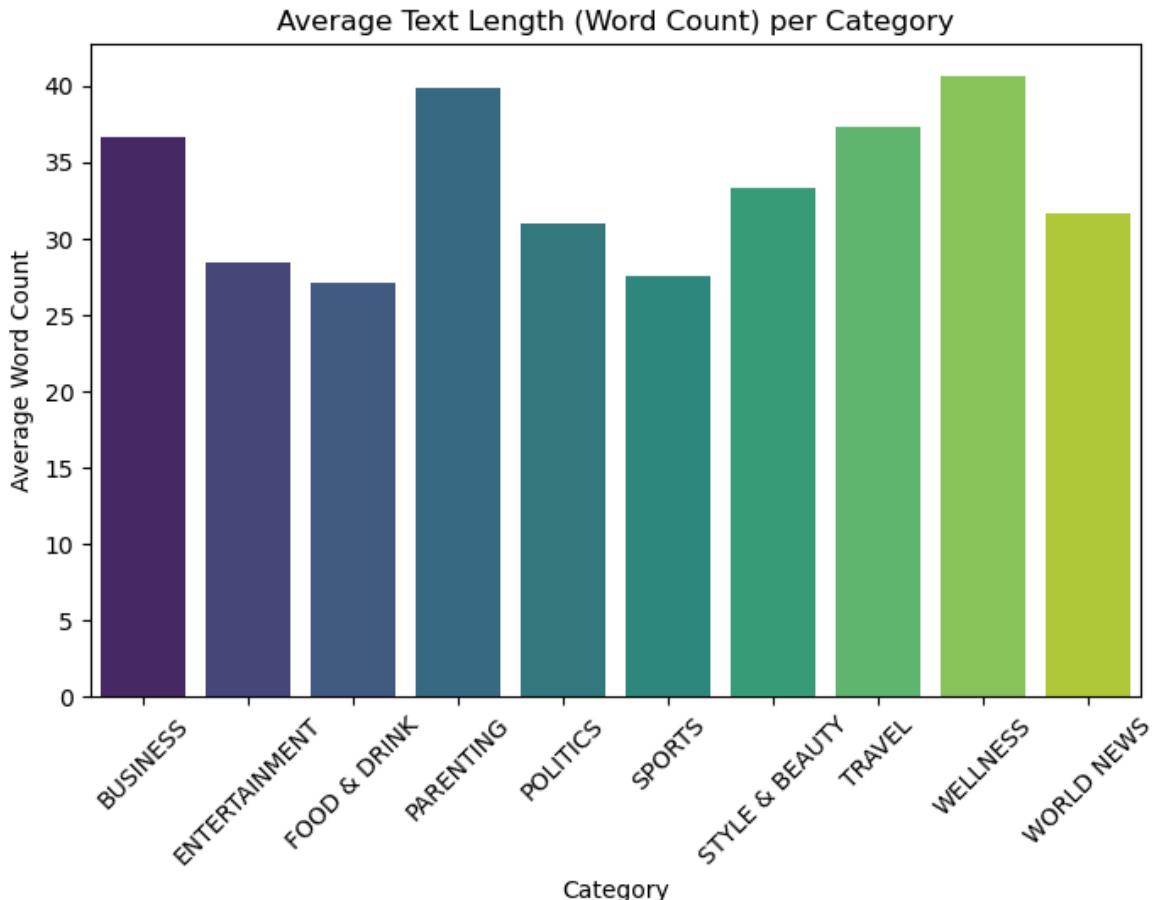
# Word and character counts on combined text
datanews_df['word_count']=datanews_df['combine_text'].apply(lambda x:len(str(x)))
datanews_df['char_count']=datanews_df['combine_text'].apply(lambda x:len(str(x)))
```

```
In [20]: # Group by category and calculate mean word count
avg_length = datanews_df.groupby('category')['word_count'].mean().reset_index()
print("Average Length =",avg_length)

#visualization (bar chart)

plt.figure(figsize=(8,5))
sns.barplot(x='category', y='word_count', data=avg_length, palette='viridis')
plt.xticks(rotation=45)
plt.title("Average Text Length (Word Count) per Category")
plt.ylabel("Average Word Count")
plt.xlabel("Category")
plt.show()
```

	category	word_count
0	BUSINESS	36.639599
1	ENTERTAINMENT	28.425000
2	FOOD & DRINK	27.130400
3	PARENTING	39.859800
4	POLITICS	30.972000
5	SPORTS	27.545147
6	STYLE & BEAUTY	33.298600
7	TRAVEL	37.325000
8	WELLNESS	40.697200
9	WORLD NEWS	31.680000

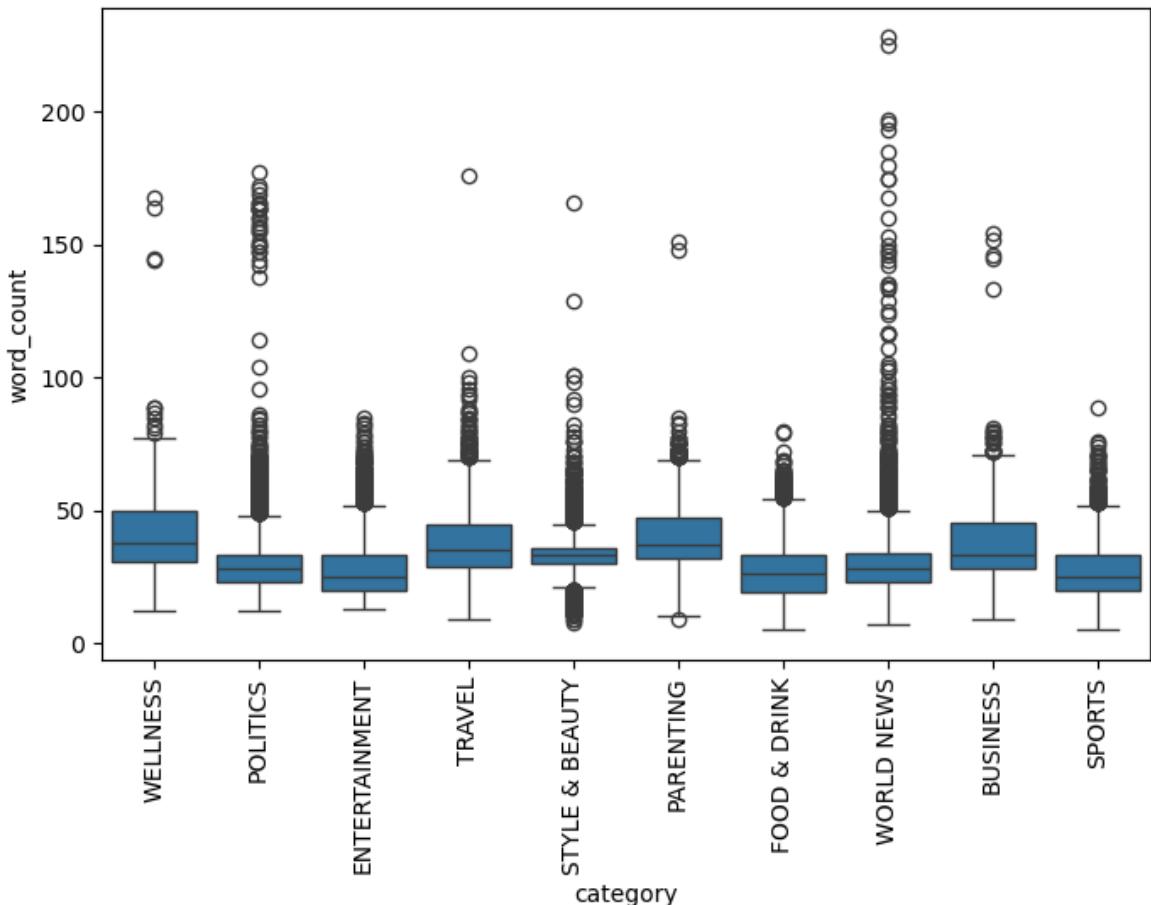


WELLNESS and PARENTING articles are longest on average; FOOD & DRINK has the shortest content.

This suggests deeper storytelling in some categories and concise formats in others.

Outlier Detection

```
In [21]: from pandas.plotting import parallel_coordinates
plt.figure(figsize=(8, 5))
sns.boxplot(x='category', y='word_count', data=datanews_df)
plt.xticks(rotation=90)
plt.show()
```



STYLE & BEAUTY and **WELLNESS** categories show the widest spread and most outliers in word count.

These outliers may reflect unusually long or short articles that could skew model training.

```
In [22]: Q1, Q3 = datanews_df['word_count'].quantile([0.25, 0.75]);
IQR = Q3 - Q1
outliers = datanews_df[(datanews_df['word_count'] < Q1 - 1.5*IQR) | (datanews_df
outliers.head())]
```

Out[22]:

	category	headline	links	short_descrip
10	WELLNESS	Conflict: 7 Decisions to Learn Instead of Fight	https://www.huffingtonpost.com/entry/conflict-...	If the o person n opens to ca cc
46	WELLNESS	Tis the Season to Be Stressed: A Survival Guide	https://www.huffingtonpost.com/entry/holiday-s...	This week, harried prep Thanksgiving
124	WELLNESS	5 Workout Wardrobe Fails and How to Avoid Them	https://www.huffingtonpost.com/entry/workout-t...	I've had earl that malfunct sports bi
141	WELLNESS	Dating Fears: 'I Don't Know The First Thing A...	https://www.huffingtonpost.com/entry/dating-fe...	I don't know first thing ab choosing
165	WELLNESS	3 Ways You Zap Happiness (And How to Get It Back)	https://www.huffingtonpost.com/entry/happiness...	Happiness i approach to It's a min



Text Normalization

In [23]:

```
#convert into Lowercase
datanews_df['combine_text']=datanews_df['combine_text'].str.lower()
datanews_df['category']=datanews_df['category'].str.lower()
datanews_df[['combine_text','category']].head()
```

Out[23]:

	combine_text	category
0	143 miles in 35 days: lessons learned resting ...	wellness
1	talking to yourself: crazy or crazy helpful? t...	wellness
2	crenezumab: trial will gauge whether alzheimer...	wellness
3	oh, what a difference she made if you want to ...	wellness
4	green superfoods first, the bad news: soda bre...	wellness

```
In [24]: #Remove puctuations
import string
datanews_df['combine_text'] = datanews_df['combine_text'].str.translate(str.maketrans('', '', string.punctuation))
datanews_df[['combine_text','category']].head()
```

Out[24]:

	combine_text	category
0	143 miles in 35 days lessons learned resting i...	wellness
1	talking to yourself crazy or crazy helpful thi...	wellness
2	crenezumab trial will gauge whether alzheimers...	wellness
3	oh what a difference she made if you want to b...	wellness
4	green superfoods first the bad news soda bread...	wellness

```
In [25]: # Remove numbers
datanews_df['combine_text'] = datanews_df['combine_text'].str.replace(r'\d+', ' ')
datanews_df[['combine_text']].head()
```

Out[25]:

	combine_text
0	miles in days lessons learned resting is par...
1	talking to yourself crazy or crazy helpful thi...
2	crenezumab trial will gauge whether alzheimers...
3	oh what a difference she made if you want to b...
4	green superfoods first the bad news soda bread...

```
In [26]: # Remove Leading, trailing, and multiple spaces
datanews_df['combine_text'] = datanews_df['combine_text'].str.strip().str.replace('  ', ' ')
datanews_df[['combine_text']].head()
```

Out[26]:

	combine_text
0	miles in days lessons learned resting is part ...
1	talking to yourself crazy or crazy helpful thi...
2	crenezumab trial will gauge whether alzheimers...
3	oh what a difference she made if you want to b...
4	green superfoods first the bad news soda bread...

Remove Stopwords

```
In [27]: # Remove stopwords
from nltk.corpus import stopwords

stop_words=set(stopwords.words('english'))
datanews_df['combine_text'] = datanews_df['combine_text'].apply(lambda x: ' '.join(x.split() - stop_words))

datanews_df[['combine_text']].head()
```

Out[27]:

	combine_text
0	miles days lessons learned resting part traini...
1	talking crazy crazy helpful think talking tool...
2	crenezumab trial gauge whether alzheimers drug...
3	oh differ made want busi keep tri perfect want...
4	green superfoods first bad news soda bread cor...

Stemming and Lemmatizing

In [28]:

```
from nltk.stem import PorterStemmer, WordNetLemmatizer
stemmer = PorterStemmer()
lemmatizer = WordNetLemmatizer()

def normalize_text(text):
    return ' '.join([stemmer.stem(lemmatizer.lemmatize(word)) for word in text])

datanews_df['clean_text'] = datanews_df['combine_text'].apply(normalize_text)
datanews_df[['clean_text', 'category']].head()
```

Out[28]:

	clean_text	category
0	mile day lesson learn rest part train ive conf...	wellness
1	talk crazi crazi help think talk tool coach ch...	wellness
2	crenezumab trial gaug whether alzhaim drug pre...	wellness
3	oh differ made want busi keep tri perfect want...	wellness
4	green superfood first bad news soda bread corn...	wellness

Feature Extraction

In [29]:

```
# One-hot encoding for category column
from sklearn.preprocessing import OneHotEncoder

encoder = OneHotEncoder()
category_encoded = encoder.fit_transform(datanews_df[['category']])
print("One-Hot Encoded for Category is done")
print(category_encoded.shape)
```

One-Hot Encoded for Category is done
(45749, 10)

In [30]:

```
from sklearn.feature_extraction.text import TfidfVectorizer

# TF-IDF vectorization
tfidf_vectorizer = TfidfVectorizer(max_features=2000, ngram_range=(1,2))
X_tfidf = tfidf_vectorizer.fit_transform(datanews_df['clean_text'])
print("TF-IDF Vectorization is done")
print(X_tfidf.shape)
```

TF-IDF Vectorization is done
(45749, 2000)

```
In [31]: # Decide X and y (true output column)
from scipy.sparse import hstack

X = X_tfidf
y = datanews_df['category']
print("TF-IDF shape:", X.shape)
```

TF-IDF shape: (45749, 2000)

```
In [32]: !pip install gensim
```

Requirement already satisfied: gensim in c:\users\vanda\anaconda3\lib\site-packages (4.4.0)
Requirement already satisfied: numpy>=1.18.5 in c:\users\vanda\anaconda3\lib\site-packages (from gensim) (2.1.3)
Requirement already satisfied: scipy>=1.7.0 in c:\users\vanda\anaconda3\lib\site-packages (from gensim) (1.15.3)
Requirement already satisfied: smart_open>=1.8.1 in c:\users\vanda\anaconda3\lib\site-packages (from gensim) (7.5.0)
Requirement already satisfied: wrapt in c:\users\vanda\anaconda3\lib\site-packages (from smart_open>=1.8.1->gensim) (1.17.0)

Word Embedding

```
In [33]: #Word Embedding

from gensim.models import Word2Vec

sentences = [text.split() for text in datanews_df['clean_text']]

# Train Word2Vec model
w2v_model = Word2Vec(
    sentences,           # tokenized text
    vector_size=100,     # embedding dimension (you can try 50, 100, 300)
    window=5,            # context window size
    min_count=2,          # ignore rare words
    workers=4             # parallel training
)

def get_doc_embedding(text):
    words = text.split()
    vectors = [w2v_model.wv[w] for w in words if w in w2v_model.wv]
    if len(vectors) == 0:
        return np.zeros(100) # fallback if no known words
    return np.mean(vectors, axis=0)

X_emb = np.array([get_doc_embedding(t) for t in datanews_df['clean_text']])

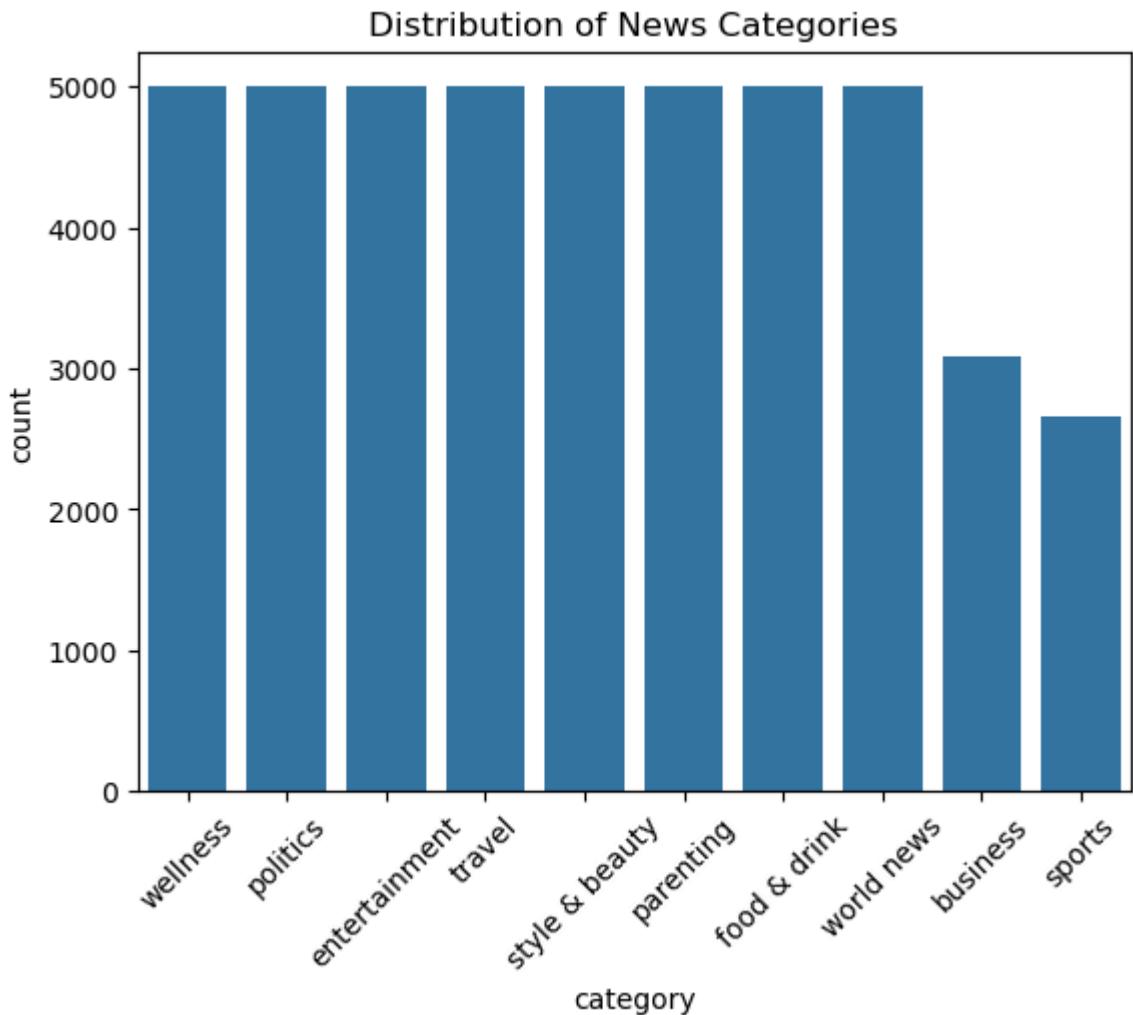
y_emb = datanews_df['category']
print("Word2Vec Embedding shape:", X_emb.shape)
```

Word2Vec Embedding shape: (45749, 100)

Category Distribution

```
In [34]: import matplotlib.pyplot as plt
import seaborn as sns

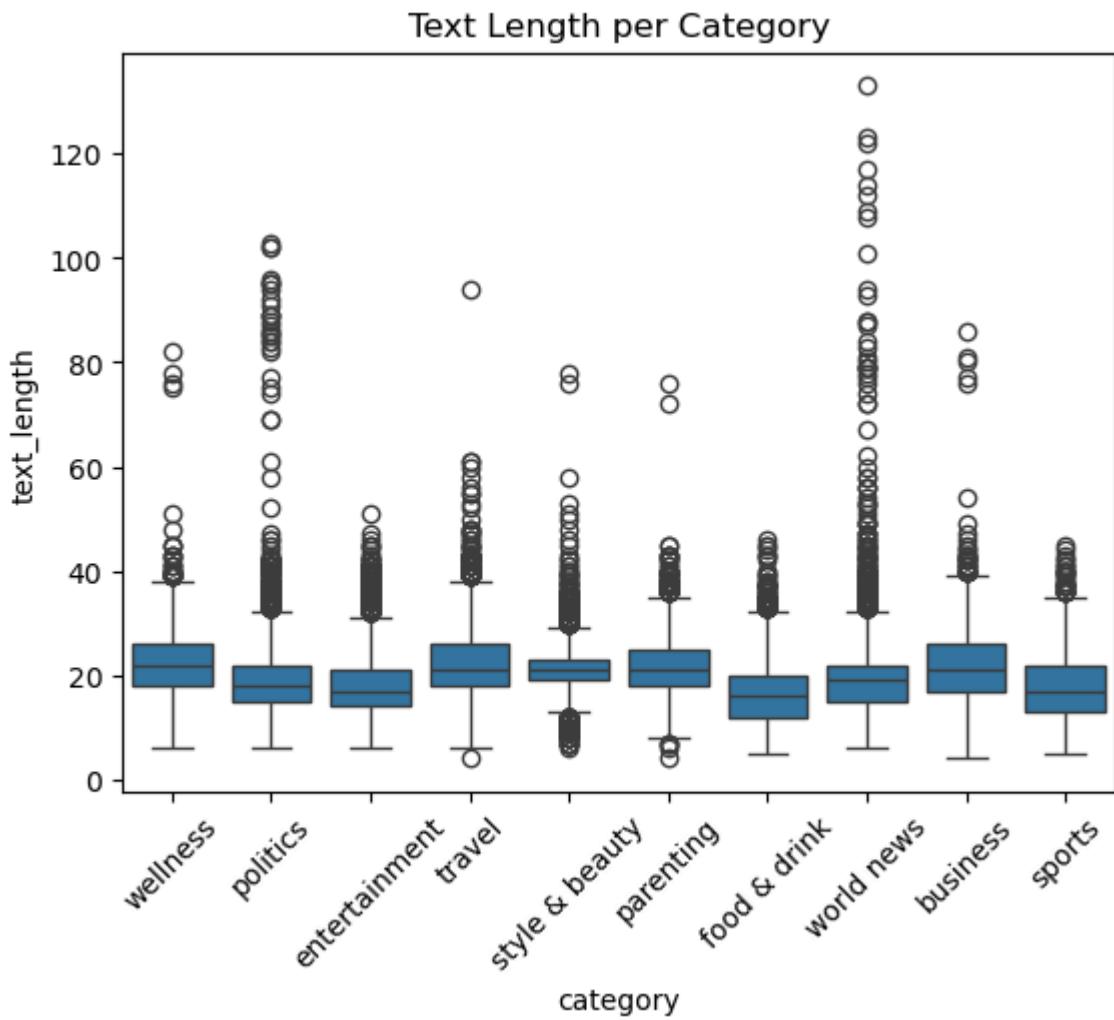
sns.countplot(x=datanews_df['category'])
plt.title("Distribution of News Categories")
plt.xticks(rotation=45)
plt.show()
```



Text Length Analysis

```
In [35]: datanews_df['text_length'] = datanews_df['clean_text'].apply(lambda x: len(x.split()))

sns.boxplot(x='category', y='text_length', data=datanews_df)
plt.title("Text Length per Category")
plt.xticks(rotation=45)
plt.show()
```



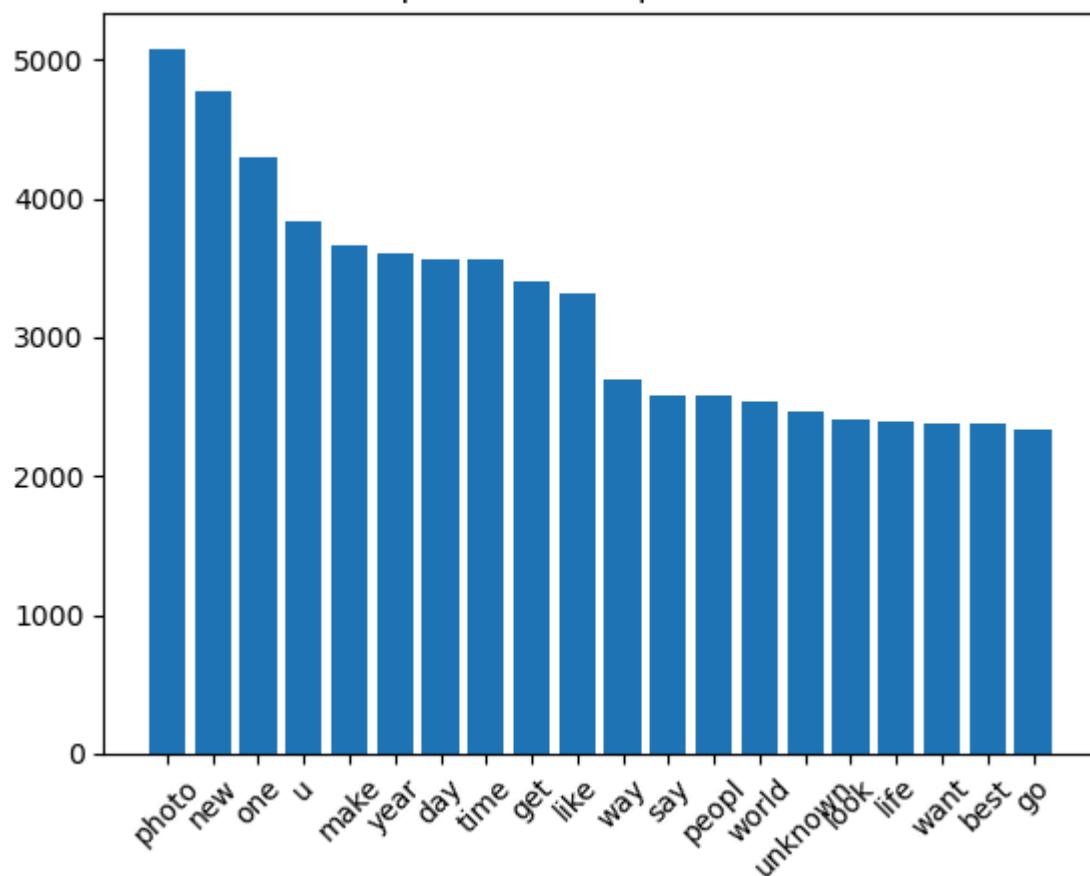
Most Frequent Words

```
In [36]: # check for most frequent words
from collections import Counter

all_words = ' '.join(datanews_df['clean_text']).split()
word_freq = Counter(all_words)
common_words = word_freq.most_common(20)

words, counts = zip(*common_words)
plt.bar(words, counts)
plt.xticks(rotation=45)
plt.title("Top 20 Most Frequent Words")
plt.show()
```

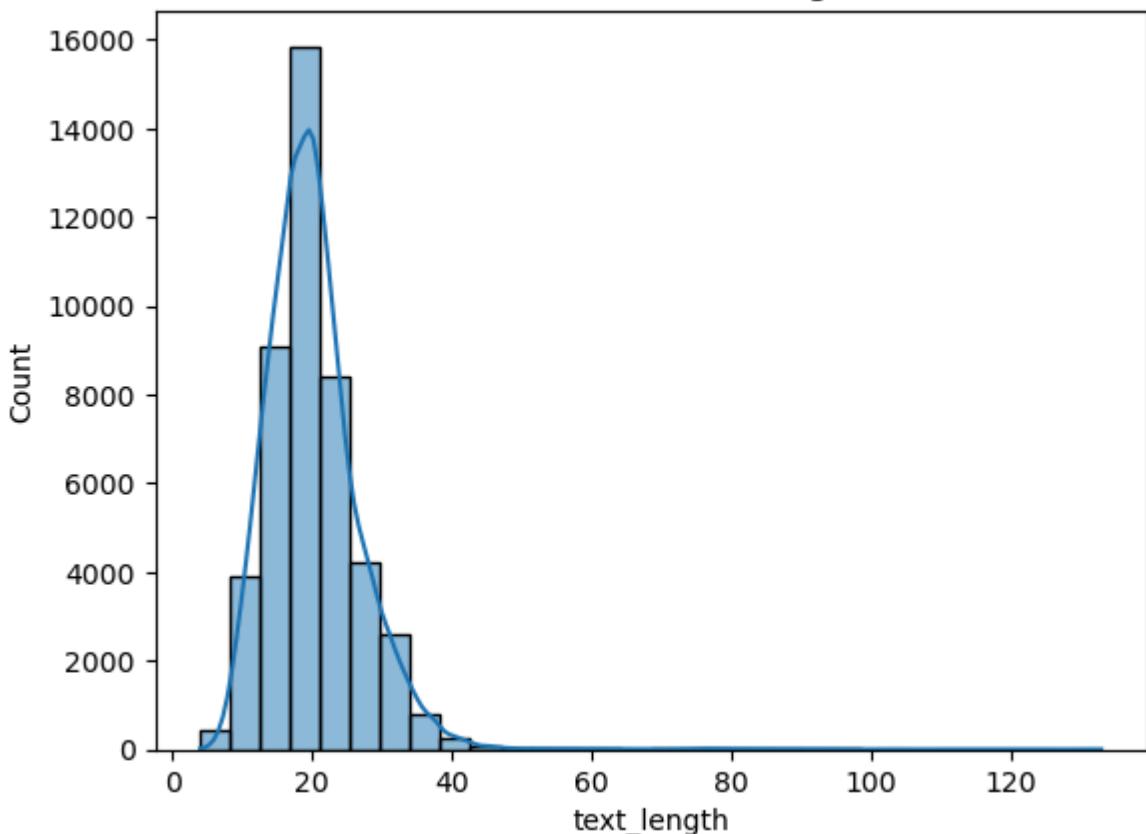
Top 20 Most Frequent Words



Compare Category vs Embedding Length

```
In [37]: sns.histplot(datanews_df['text_length'], bins=30, kde=True)
plt.title("Distribution of Text Lengths")
plt.show()
```

Distribution of Text Lengths



Data Splitting

```
In [38]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

Model Building

Logistic Regression

```
In [39]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(class_weight='balanced', max_iter=1000)
lr.fit(X_train, y_train)
```

Out[39]:

▾ LogisticRegression
 ⓘ ⓘ

 LogisticRegression(class_weight='balanced', max_iter=1000)

Naive Bayes

```
In [55]: from sklearn.naive_bayes import MultinomialNB
nb = MultinomialNB()
nb.fit(X_train, y_train)
```

```
Out[55]: ▾ MultinomialNB ⓘ ?  
MultinomialNB()
```

SVM

```
In [57]: from sklearn.svm import LinearSVC  
  
svm = LinearSVC(class_weight='balanced')  
svm.fit(X_train, y_train)
```

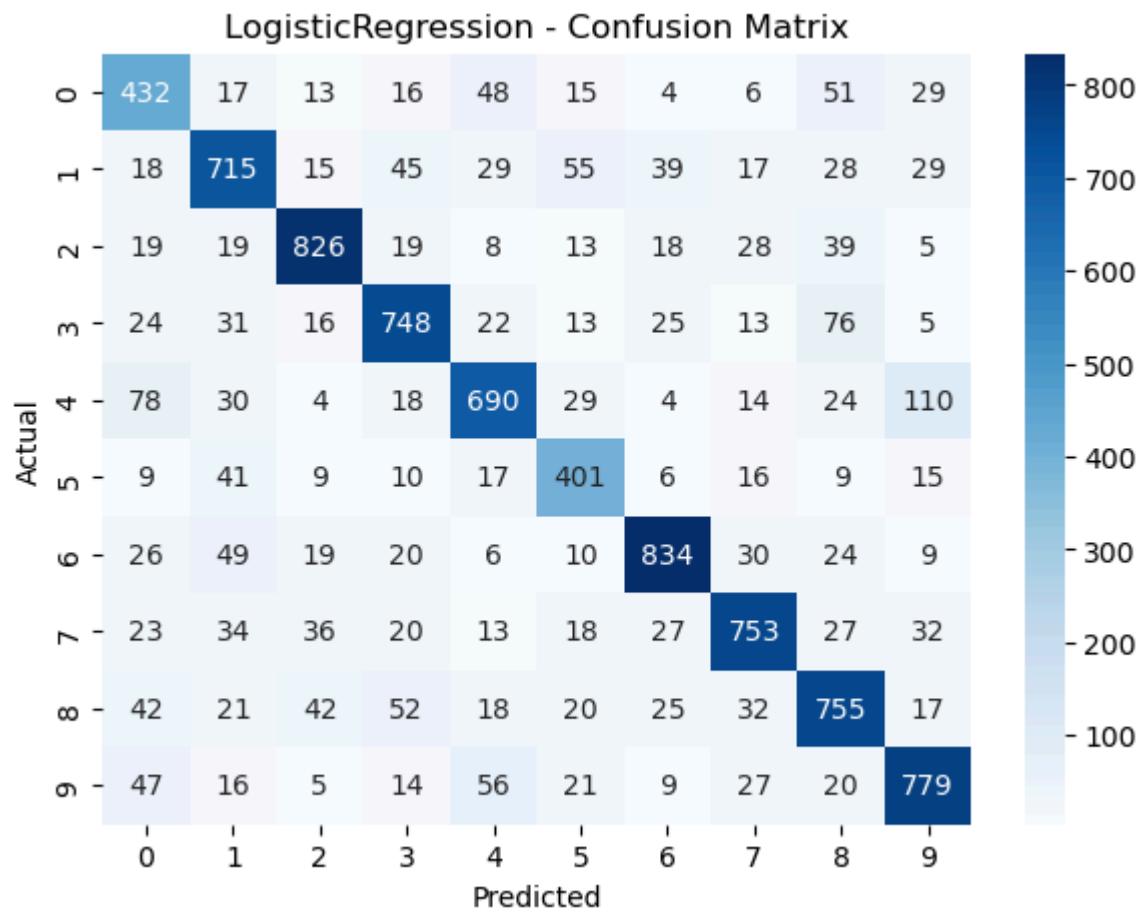
```
Out[57]: ▾ LinearSVC ⓘ ?  
LinearSVC(class_weight='balanced')
```

Model Evaluation

```
In [42]: from sklearn.metrics import classification_report,confusion_matrix  
  
# Dictionary to store accuracy  
accuracy_scores = {}  
  
for model in [lr, nb, svm]:  
    y_pred = model.predict(X_test)  
    acc = model.score(X_test, y_test)  
    accuracy_scores[model.__class__.__name__] = acc  
  
    # Print reports  
    print(f"\nModel: {model.__class__.__name__}")  
    print(classification_report(y_test, y_pred))  
    print("\nAccuracy:", acc)  
    print("\n\n")  
  
    #Confusion matrix visualization  
    plt.figure(figsize=(7,5))  
    sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blu  
    plt.xlabel("Predicted")  
    plt.ylabel("Actual")  
    plt.title(f"{model.__class__.__name__} - Confusion Matrix")  
    plt.show()
```

Model: LogisticRegression		precision	recall	f1-score	support
business	0.60	0.68	0.64	631	
entertainment	0.73	0.72	0.73	990	
food & drink	0.84	0.83	0.83	994	
parenting	0.78	0.77	0.77	973	
politics	0.76	0.69	0.72	1001	
sports	0.67	0.75	0.71	533	
style & beauty	0.84	0.81	0.83	1027	
travel	0.80	0.77	0.78	983	
wellness	0.72	0.74	0.73	1024	
world news	0.76	0.78	0.77	994	
accuracy			0.76	9150	
macro avg	0.75	0.75	0.75	9150	
weighted avg	0.76	0.76	0.76	9150	

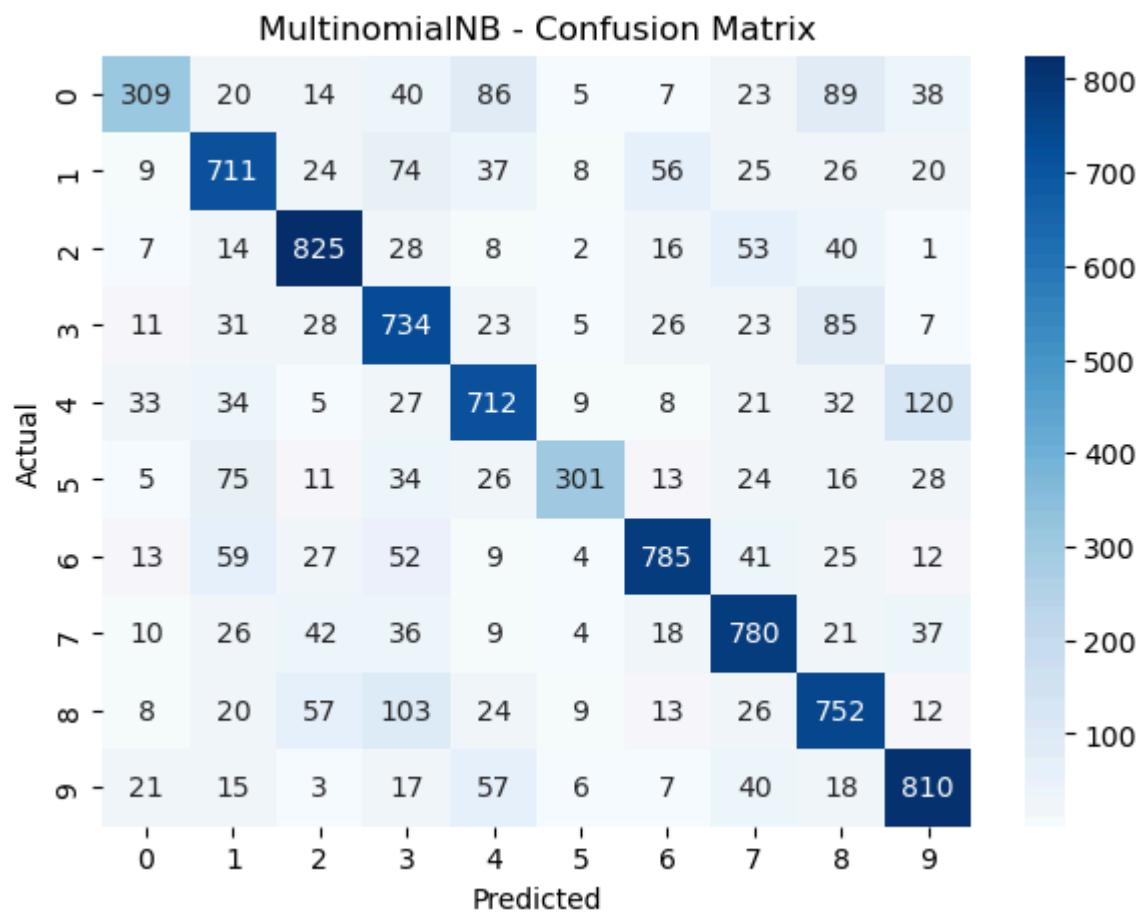
Accuracy: 0.7577049180327868



Model: MultinomialNB

	precision	recall	f1-score	support
business	0.73	0.49	0.58	631
entertainment	0.71	0.72	0.71	990
food & drink	0.80	0.83	0.81	994
parenting	0.64	0.75	0.69	973
politics	0.72	0.71	0.71	1001
sports	0.85	0.56	0.68	533
style & beauty	0.83	0.76	0.79	1027
travel	0.74	0.79	0.77	983
wellness	0.68	0.73	0.71	1024
world news	0.75	0.81	0.78	994
accuracy			0.73	9150
macro avg	0.74	0.72	0.72	9150
weighted avg	0.74	0.73	0.73	9150

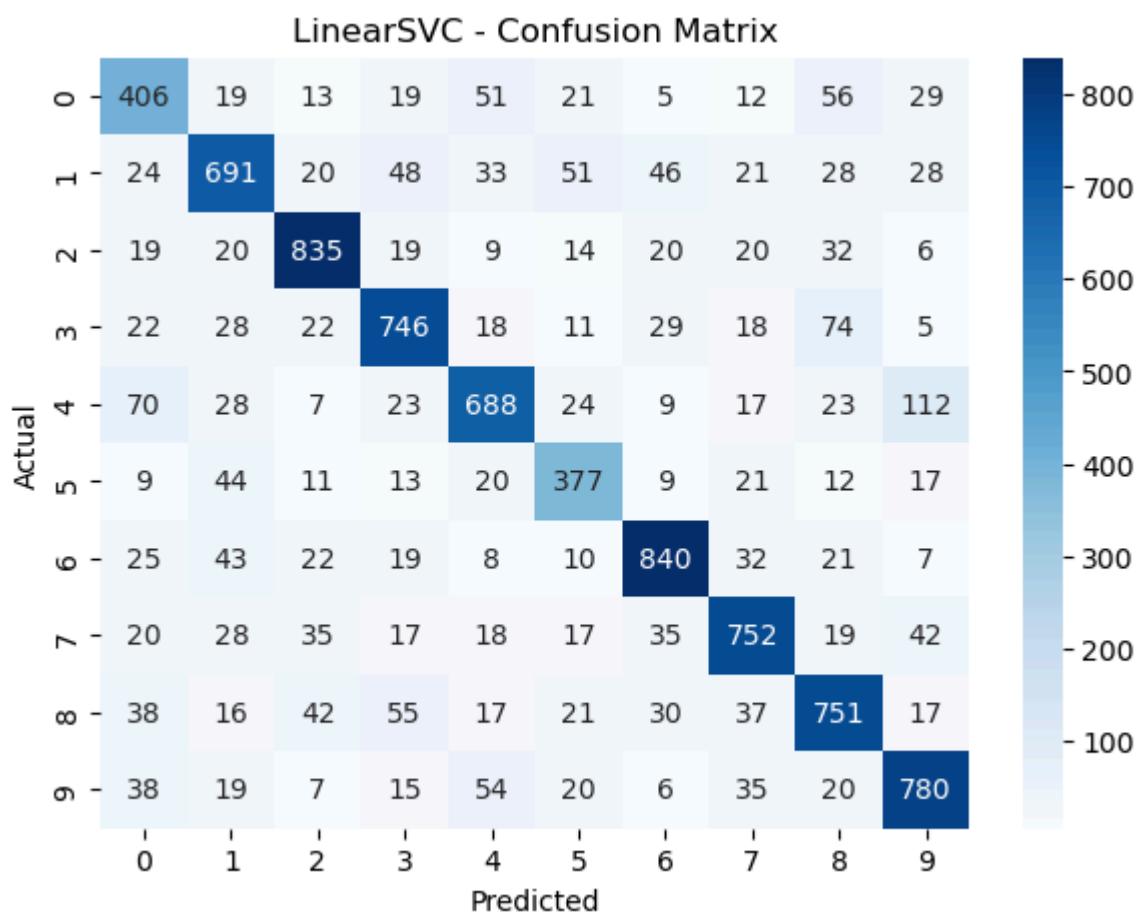
Accuracy: 0.7343169398907103



Model: LinearSVC

	precision	recall	f1-score	support
business	0.61	0.64	0.62	631
entertainment	0.74	0.70	0.72	990
food & drink	0.82	0.84	0.83	994
parenting	0.77	0.77	0.77	973
politics	0.75	0.69	0.72	1001
sports	0.67	0.71	0.69	533
style & beauty	0.82	0.82	0.82	1027
travel	0.78	0.77	0.77	983
wellness	0.72	0.73	0.73	1024
world news	0.75	0.78	0.77	994
accuracy			0.75	9150
macro avg	0.74	0.74	0.74	9150
weighted avg	0.75	0.75	0.75	9150

Accuracy: 0.7503825136612022



Check for Feature Importance

```
In [44]: # Get feature names from TF-IDF vectorizer
feature_names = tfidf_vectorizer.get_feature_names_out()

# Coefficients from the best Logistic Regression model
coefficients = best_lr.coef_
```

```
# Map each class to its top features
for idx, category in enumerate(best_lr.classes_):
    top10_idx = np.argsort(coefficients[idx])[-10:] # top 10 positive
    bottom10_idx = np.argsort(coefficients[idx])[:10] # top 10 negative

    print(f"\nCategory: {category}")
    print("Top positive words:", [feature_names[i] for i in top10_idx])
    print("Top negative words:", [feature_names[i] for i in bottom10_idx])
```

Category: business

Top positive words: ['bank', 'wall st', 'employe', 'walmart', 'uber', 'market', 'job', 'ceo', 'compani', 'busi']

Top negative words: ['photo', 'child', 'travel', 'kid', 'fashion', 'mom', 'republ ican', 'recip', 'eat', 'senat']

Category: entertainment

Top positive words: ['kardashian', 'box offic', 'album', 'music', 'singer', 'acto r', 'movi', 'jenner', 'trailer', 'film']

Top negative words: ['photo', 'food', 'recip', 'fashion', 'travel', 'style', 'unk nown', 'kid', 'parent', 'state']

Category: food & drink

Top positive words: ['cake', 'kitchen', 'tast', 'chef', 'chees', 'cook', 'wine', 'cocktail', 'food', 'recip']

Top negative words: ['travel', 'kid', 'parent', 'health', 'music', 'feel', 'calor i', 'hair', 'child', 'fashion']

Category: parenting

Top positive words: ['pregnanc', 'son', 'breastfeed', 'toddler', 'daughter', 'bab i', 'child', 'mom', 'parent', 'kid']

Top negative words: ['travel', 'recip', 'fashion', 'welcom', 'photo', 'hotel', 's tyle', 'singer', 'jame', 'citi']

Category: politics

Top positive words: ['gun', 'congress', 'obama', 'rubio', 'immigr', 'governor', 'democrat', 'senat', 'republican', 'gop']

Top negative words: ['photo', 'best', 'parent', 'travel', 'fashion', 'recip', 'ki d', 'movi', 'style', 'season']

Category: sports

Top positive words: ['team', 'win', 'coach', 'game', 'olymp', 'footbal', 'sport', 'player', 'nba', 'nfl']

Top negative words: ['photo', 'food', 'recip', 'travel', 'parent', 'movi', 'fashi on', 'music', 'film', 'child']

Category: style & beauty

Top positive words: ['beauti', 'kate', 'wear', 'makeup', 'hair', 'dress', 'mode l', 'style', 'photo', 'fashion']

Top negative words: ['recip', 'food', 'hotel', 'kid', 'travel', 'parent', 'worl d', 'eat', 'war', 'child']

Category: travel

Top positive words: ['place', 'getaway', 'trip', 'destin', 'flight', 'tsa', 'vaca t', 'airlin', 'hotel', 'travel']

Top negative words: ['recip', 'mom', 'attack', 'parent', 'style', 'health', 'chil d', 'fashion', 'kill', 'babi']

Category: wellness

Top positive words: ['weight', 'healthi', 'yoga', 'exercis', 'research', 'cance r', 'medit', 'sleep', 'health', 'workout']

Top negative words: ['photo', 'kid', 'parent', 'child', 'travel', 'babi', 'mom', 'recip', 'fashion', 'busi']

Category: world news

Top positive words: ['israel', 'migrant', 'syria', 'syrian', 'kill', 'saudi', 'isi', 'turkey', 'attack', 'refugee']

Top negative words: ['photo', 'travel', 'star', 'republican', 'best', 'love', 'ki d', 'gop', 'think', 'recip']

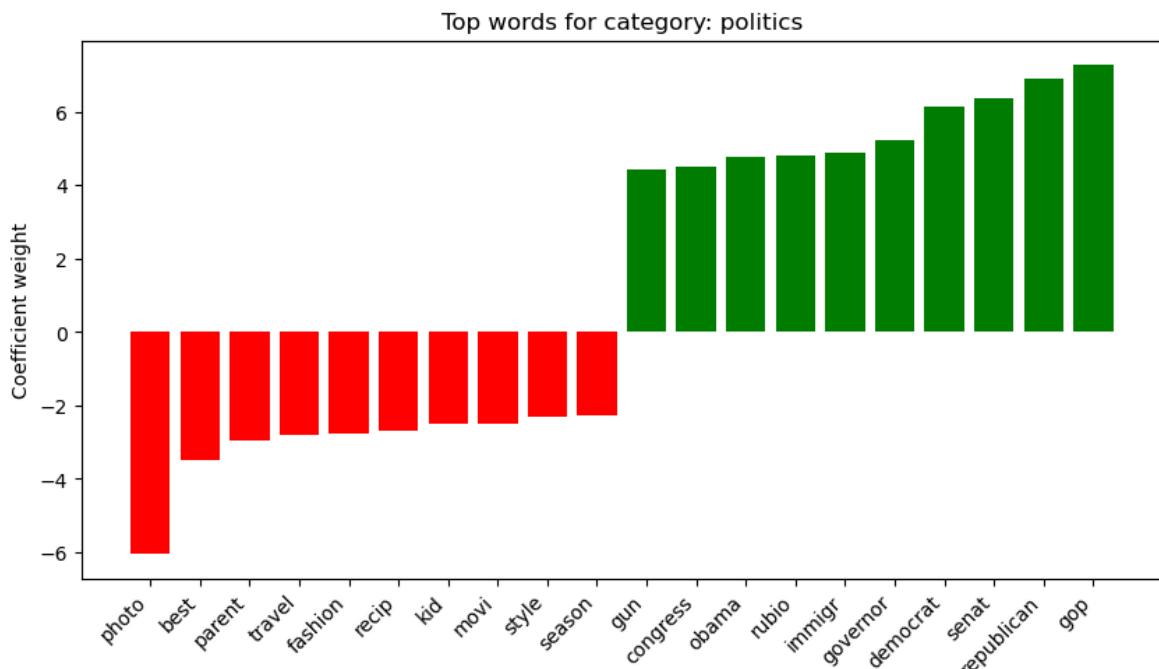
```
In [45]: def plot_top_words(classifier, feature_names, category, class_index, top_n=10):
    # Get coefficients for the category
    coef = classifier.coef_[class_index]

    # Top positive and negative indices
    top_pos_idx = np.argsort(coef)[-top_n:]
    top_neg_idx = np.argsort(coef)[:top_n]

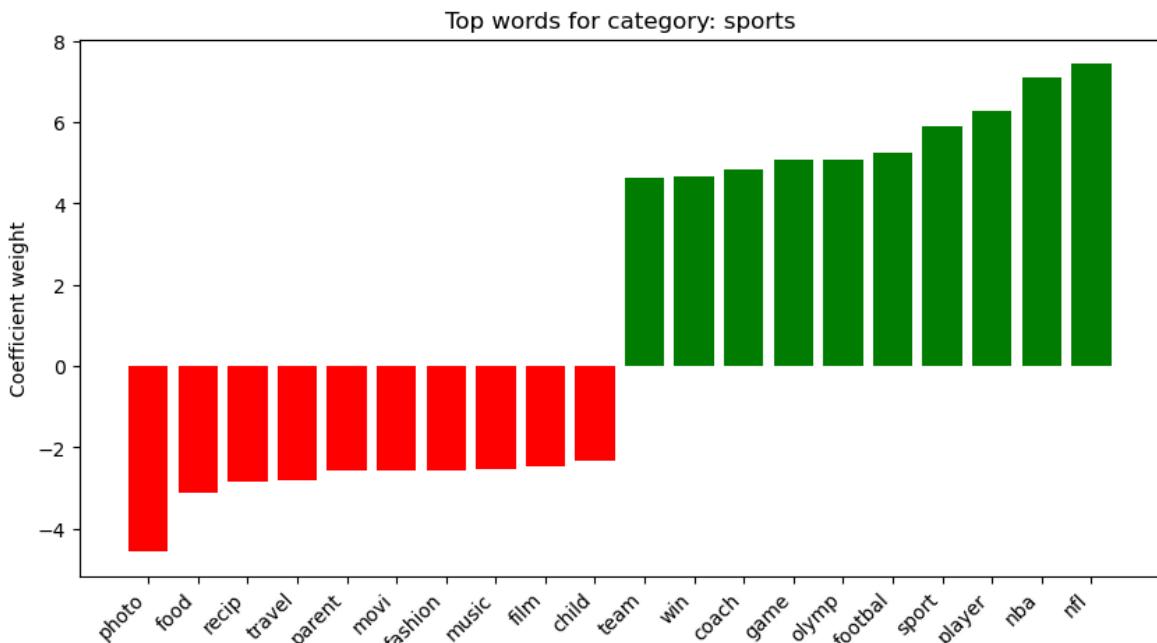
    # Combine indices
    top_idx = np.concatenate([top_neg_idx, top_pos_idx])

    # Plot
    plt.figure(figsize=(10,5))
    colors = ['red']*top_n + ['green']*top_n
    plt.bar(range(2*top_n), coef[top_idx], color=colors)
    plt.xticks(range(2*top_n), [feature_names[i] for i in top_idx], rotation=45,
               fontstyle='italic')
    plt.title(f"Top words for category: {category}")
    plt.ylabel("Coefficient weight")
    plt.show()

    # Example: plot for 'politics' category
plot_top_words(best_lr, feature_names, "politics", list(best_lr.classes_).index("s"))
```



```
In [46]: # Example: plot for 'politics' category
plot_top_words(best_lr, feature_names, "sports", list(best_lr.classes_).index("s"))
```



Hyperparameter Tuning

Logistic Regression with Hyper parameter Tuning — RandomizedSearchCV

```
In [43]: from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import classification_report, accuracy_score

# Define model
lr = LogisticRegression(multi_class='auto')

# Define hyperparameters
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'max_iter': [100, 200, 500, 1000],
    'solver': ['liblinear', 'lbfgs']
}

# RandomizedSearchCV
random_search_lr = RandomizedSearchCV(
    estimator=lr,
    param_distributions=param_grid,
    n_iter=10,
    cv=5,
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)

# Fit on training data
random_search_lr.fit(X_train, y_train)

# Get best model
best_lr = random_search_lr.best_estimator_

# Evaluate best model
```

```
y_pred = best_lr.predict(X_test)

print("Best Parameters:", random_search_lr.best_params_)
print("Best CV Accuracy:", random_search_lr.best_score_)
print("Test Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Best Parameters: {'solver': 'liblinear', 'max_iter': 200, 'C': 1}
 Best CV Accuracy: 0.7503211045135163
 Test Accuracy: 0.7572677595628415

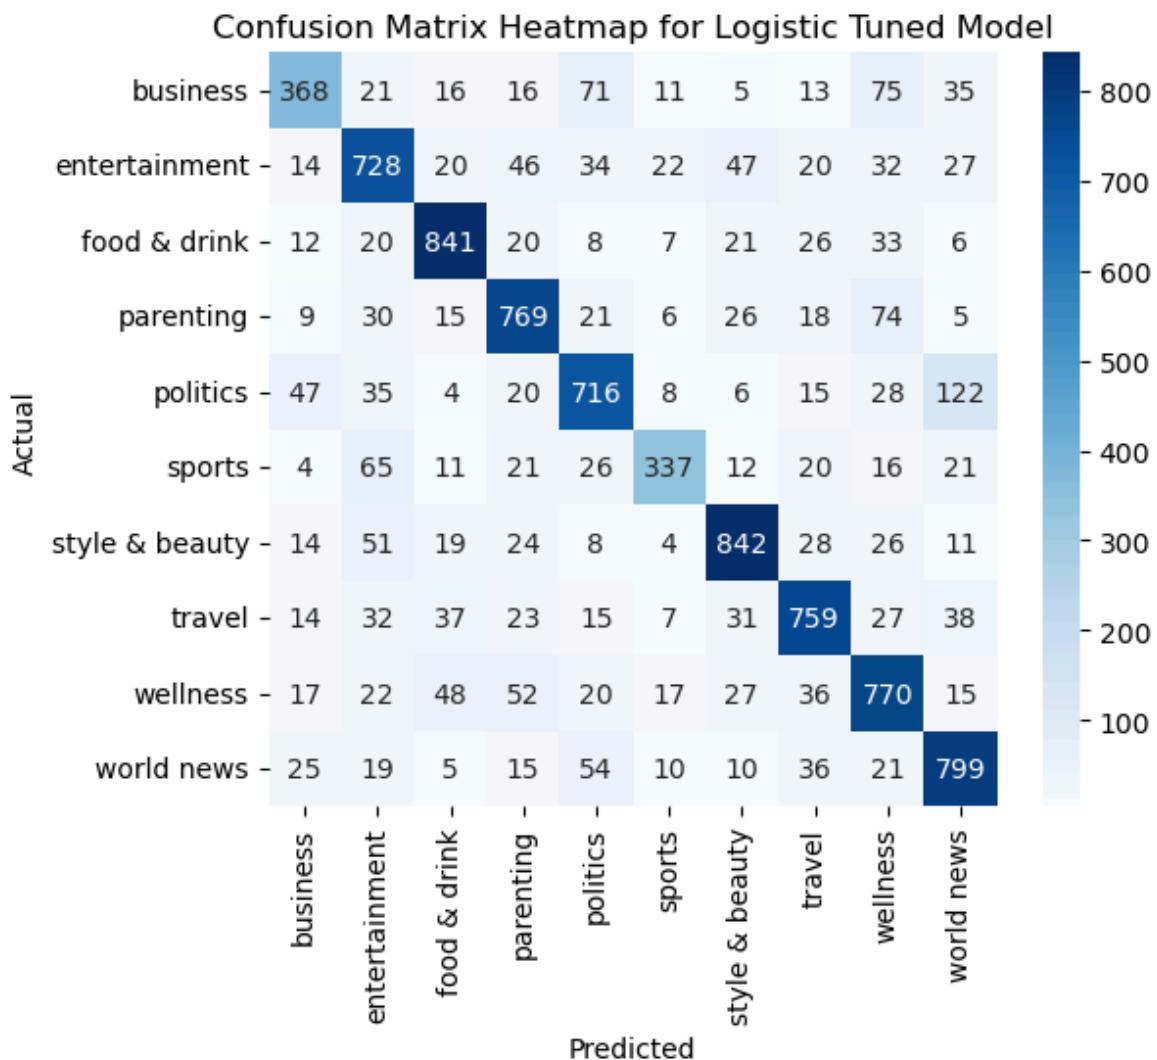
Classification Report:

	precision	recall	f1-score	support
business	0.70	0.58	0.64	631
entertainment	0.71	0.74	0.72	990
food & drink	0.83	0.85	0.84	994
parenting	0.76	0.79	0.78	973
politics	0.74	0.72	0.73	1001
sports	0.79	0.63	0.70	533
style & beauty	0.82	0.82	0.82	1027
travel	0.78	0.77	0.78	983
wellness	0.70	0.75	0.72	1024
world news	0.74	0.80	0.77	994
accuracy			0.76	9150
macro avg	0.76	0.75	0.75	9150
weighted avg	0.76	0.76	0.76	9150

Confusion Matrix Heatmap for Logistic Regression

```
In [47]: # Generate confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=best_lr.classes_)

plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=best_lr.classes_,
            yticklabels=best_lr.classes_)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix Heatmap for Logistic Tuned Model")
plt.show()
```



Navie Bayes with Hyperparameter Tuning

```
In [48]: # Define model
nb = MultinomialNB()

# Define hyperparameters
param_grid = {
    'alpha': [0.1, 0.5, 1.0, 1.5, 2.0] # smoothing parameter
}

# RandomizedSearchCV
random_search_nb = RandomizedSearchCV(
    estimator=nb,
    param_distributions=param_grid,
    n_iter=5, # number of random samples from param_grid
    cv=5, # 5-fold cross-validation
    scoring='accuracy',
    random_state=42,
    n_jobs=-1
)

# Fit on training data
random_search_nb.fit(X_train, y_train)

# Get best model
best_nb = random_search_nb.best_estimator_
```

```
# Evaluate best model
y_pred_nb = best_nb.predict(X_test)

print("Best Parameters:", random_search_nb.best_params_)
print("Best CV Accuracy:", random_search_nb.best_score_)
print("Test Accuracy:", accuracy_score(y_test, y_pred_nb))
print("\nClassification Report:\n", classification_report(y_test, y_pred_nb))

Best Parameters: {'alpha': 1.0}
Best CV Accuracy: 0.7297193695277728
Test Accuracy: 0.7343169398907103

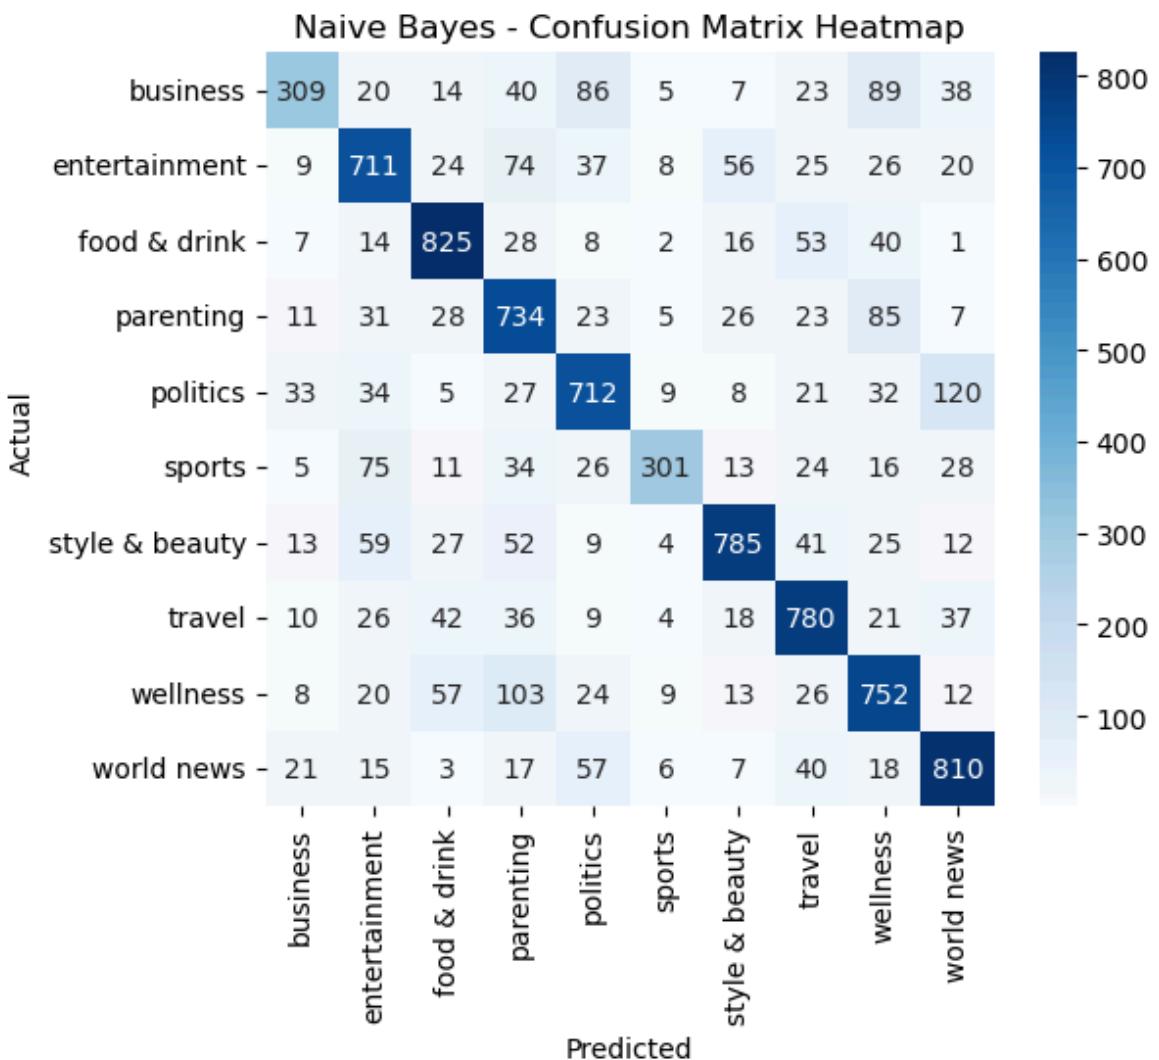
Classification Report:
precision    recall    f1-score   support
business      0.73     0.49     0.58      631
entertainment  0.71     0.72     0.71      990
food & drink  0.80     0.83     0.81      994
parenting      0.64     0.75     0.69      973
politics       0.72     0.71     0.71     1001
sports         0.85     0.56     0.68      533
style & beauty 0.83     0.76     0.79     1027
travel          0.74     0.79     0.77      983
wellness        0.68     0.73     0.71     1024
world news     0.75     0.81     0.78      994

accuracy       None      None     0.73     9150
macro avg      0.74     0.72     0.72     9150
weighted avg   0.74     0.73     0.73     9150
```

Confusion Matrix Heatmap for Naive Bayes

```
In [49]: # Generate confusion matrix
cm = confusion_matrix(y_test, y_pred_nb, labels=best_nb.classes_)

# Plot heatmap
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=best_nb.classes_,
            yticklabels=best_nb.classes_)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Naive Bayes - Confusion Matrix Heatmap")
plt.show()
```



SVM with Hyperparameter Tuning

```
In [50]: from sklearn.svm import LinearSVC
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.metrics import classification_report, accuracy_score

# Define the search space
param_grid = {
    "C": [0.1, 1, 10] # Regularization strength
}

# Set up scoring
scoring = 'accuracy'

# Set up k-fold cross-validation
kfold = StratifiedKFold(n_splits=3, shuffle=True, random_state=0)

# Define model
svm = LinearSVC()

# GridSearchCV
grid_search_svm = GridSearchCV(
    estimator=svm,
    param_grid=param_grid,
    scoring=scoring,
    cv=kfold,
```

```

    n_jobs=-1
)

# Fit on training data
grid_search_svm.fit(X_train, y_train)

# Get best model
best_svm = grid_search_svm.best_estimator_

# Evaluate best model
y_pred_svm = best_svm.predict(X_test)

print("Best Parameters:", grid_search_svm.best_params_)
print("Best CV Accuracy:", grid_search_svm.best_score_)
print("Test Accuracy:", accuracy_score(y_test, y_pred_svm))
print("\nClassification Report:\n", classification_report(y_test, y_pred_svm))

```

Best Parameters: {'C': 0.1}
 Best CV Accuracy: 0.7504850281107874
 Test Accuracy: 0.7558469945355192

Classification Report:

	precision	recall	f1-score	support
business	0.69	0.57	0.63	631
entertainment	0.72	0.73	0.73	990
food & drink	0.82	0.85	0.84	994
parenting	0.76	0.79	0.77	973
politics	0.74	0.71	0.73	1001
sports	0.76	0.64	0.70	533
style & beauty	0.81	0.82	0.82	1027
travel	0.78	0.77	0.77	983
wellness	0.70	0.75	0.73	1024
world news	0.74	0.81	0.77	994
accuracy			0.76	9150
macro avg	0.75	0.74	0.75	9150
weighted avg	0.76	0.76	0.75	9150

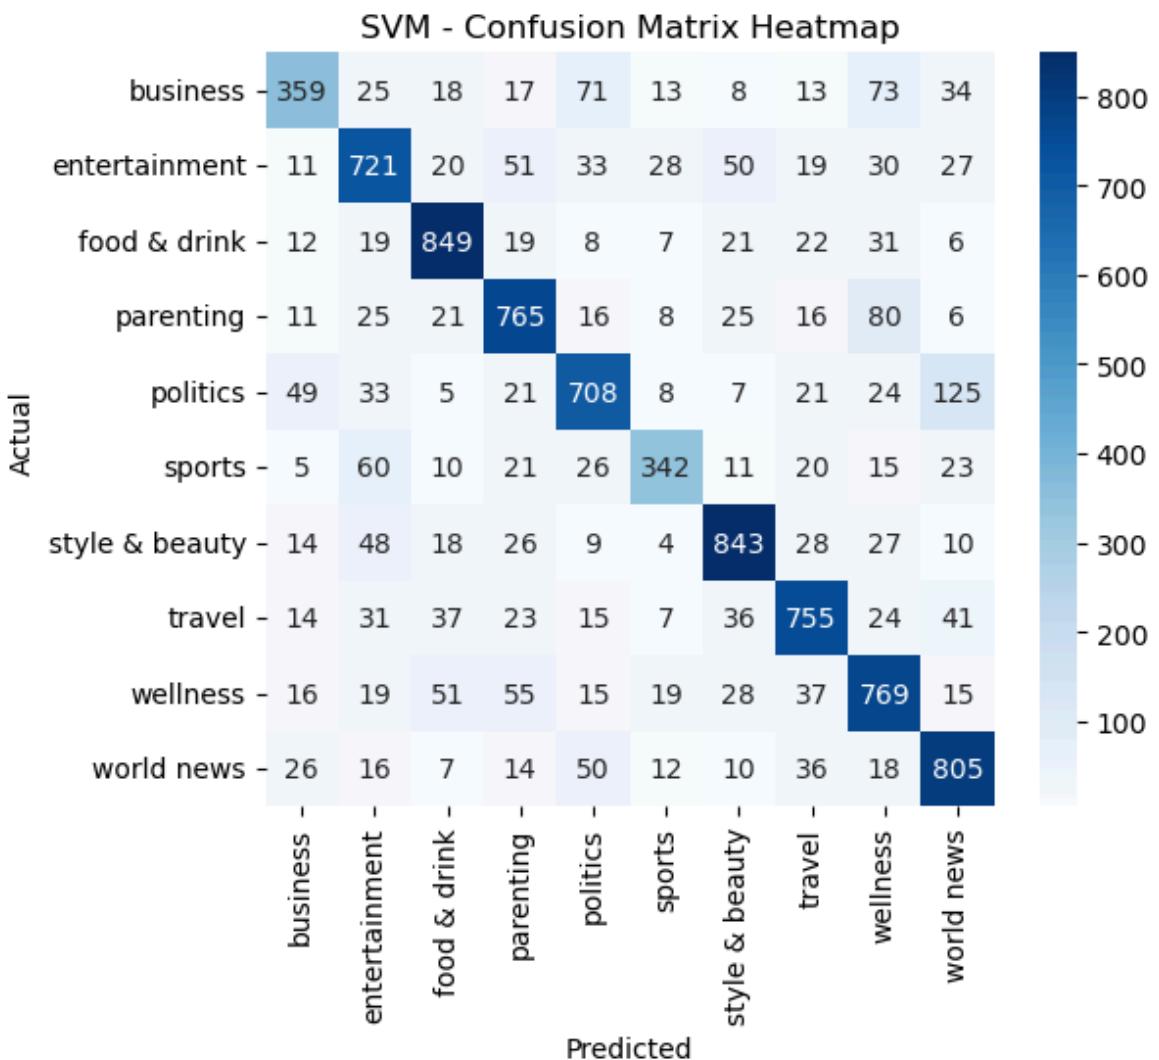
Confusion Matrix Heatmap for SVM

```

In [51]: # Generate confusion matrix
cm = confusion_matrix(y_test, y_pred_svm, labels=best_svm.classes_)

# Plot heatmap
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=best_svm.classes_,
            yticklabels=best_svm.classes_)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("SVM - Confusion Matrix Heatmap")
plt.show()

```



```
In [52]: # Model names
models = ["Logistic Regression", "Naive Bayes", "SVM (Linear)"]

# Metrics
test_accuracy = [0.7573, 0.7343, 0.7558]
macro_f1 = [0.75, 0.73, 0.75]
weighted_f1 = [0.76, 0.73, 0.75]

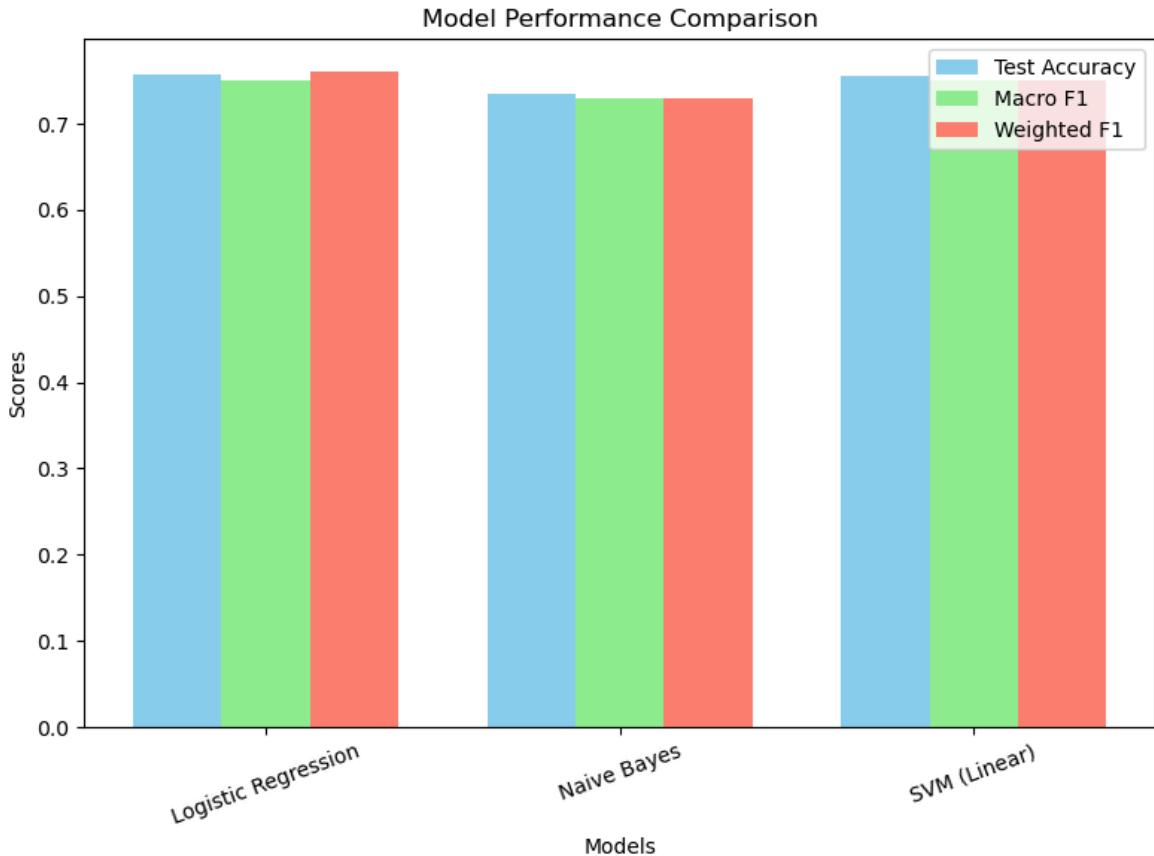
# Grouped bar chart setup
x = np.arange(len(models)) # positions
width = 0.25 # bar width

fig, ax = plt.subplots(figsize=(8,6))

# Plot bars
ax.bar(x - width, test_accuracy, width, label='Test Accuracy', color='skyblue')
ax.bar(x, macro_f1, width, label='Macro F1', color='lightgreen')
ax.bar(x + width, weighted_f1, width, label='Weighted F1', color='salmon')

# Labels and formatting
ax.set_xlabel("Models")
ax.set_ylabel("Scores")
ax.set_title("Model Performance Comparison")
ax.set_xticks(x)
ax.set_xticklabels(models, rotation=20)
ax.legend()
```

```
# Show chart
plt.tight_layout()
plt.show()
```



In [53]: `!pip install prettytable`

```
Requirement already satisfied: prettytable in c:\users\vanda\anaconda3\lib\site-packages (3.17.0)
Requirement already satisfied: wcwidth in c:\users\vanda\anaconda3\lib\site-packages (from prettytable) (0.2.5)
```

In [58]: `from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from prettytable import PrettyTable`

```
# Initialize PrettyTable with column names
table = PrettyTable()
table.field_names = ["Model", "Accuracy", "Precision", "Recall", "F1-Score"]

# List of models to compare
models = {
    "Logistic Regression": best_lr,
    "Naive Bayes": nb,
    "SVM": svm
}

# Evaluate each model and add rows to the table
for name, model in models.items():
    y_pred = model.predict(X_test)
    accuracy = round(accuracy_score(y_test, y_pred), 4)
    precision = round(precision_score(y_test, y_pred, average='macro'), 4)
    recall = round(recall_score(y_test, y_pred, average='macro'), 4)
    f1 = round(f1_score(y_test, y_pred, average='macro'), 4)

    table.add_row([name, accuracy, precision, recall, f1])
```

```
# Print the table
print("\n\033[1mModel Comparison Table\033[0m")
print(table)
```

Model Comparison Table

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.7573	0.7568	0.745	0.7493
Naive Bayes	0.7343	0.7435	0.7175	0.7243
SVM	0.7504	0.7418	0.7444	0.7427

Compare Models with Cross-Validation

```
In [59]: from sklearn.model_selection import cross_val_score
models = {
    "Logistic Regression": best_lr,
    "Naive Bayes": nb,
    "SVM": svm
}

for name, model in models.items():
    scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')
    print(f"\n{name}")
    print("CV Scores:", scores)
    print("Mean Accuracy:", scores.mean())
```

Logistic Regression

CV Scores: [0.75595628 0.75639344 0.74568306 0.7568306 0.75582031]

Mean Accuracy: 0.75413673924171

Naive Bayes

CV Scores: [0.73540984 0.73125683 0.72797814 0.73759563 0.73691114]

Mean Accuracy: 0.7338303149975482

SVM

CV Scores: [0.75147541 0.74557377 0.73956284 0.74808743 0.74707618]

Mean Accuracy: 0.7463551273482664

Model Comparison Insights

Logistic Regression

- Most balanced model across all categories.
- Fast to train and easy to interpret.
- Strong performance on majority and minority classes.
- Best choice for production and reporting.

Naive Bayes

- Fastest to train.
- Performs well on frequent categories.
- Underperforms on minority classes like 'business' and 'parenting'.

- Good for baseline or ensemble use.

SVM (Linear)

- Matches Logistic Regression in accuracy and F1-score.
- Much slower to train, especially with large TF-IDF matrices.
- Not ideal for iterative tuning or large-scale deployment.

Final Recommendation:

- Use Logistic Regression as your primary model — it offers the best trade-off between performance, speed, and interpretability.

Model Deployment

```
In [60]: import joblib
joblib.dump(best_lr, "category_model.pkl")
joblib.dump(tfidf_vectorizer, "tfidf.pkl")
```

```
Out[60]: ['tfidf.pkl']
```

```
In [61]: import joblib
from sklearn.feature_extraction.text import TfidfVectorizer

# 1. Load trained model and vectorizer
model = joblib.load("category_model.pkl")
vectorizer = joblib.load("tfidf.pkl")

# 2. Function to predict category with confidence
def predict_category(article_text):
    # Transform article into TF-IDF features
    X_input = vectorizer.transform([article_text])

    # Predict category
    prediction = model.predict(X_input)[0]

    # Confidence scores (probabilities)
    if hasattr(model, "predict_proba"):
        probs = model.predict_proba(X_input)[0]
        confidence = max(probs)
        # Map classes to probabilities
        class_probs = dict(zip(model.classes_, probs))
    else:
        # Fallback if model doesn't support predict_proba
        confidence = None
        class_probs = None

    return prediction, confidence, class_probs

# 3. Example usage
article = """
The new Michelin-starred restaurant in Mumbai is redefining Indian cuisine with
Their signature dish, saffron-infused lamb biryani, has become a favorite among
"""
predicted_category, confidence, class_probs = predict_category(article)
```

```
print("Article: \n",article)
print("Predicted Category:", predicted_category)
print("\nConfidence Score:", confidence)
```

Article:

The new Michelin-starred restaurant in Mumbai is redefining Indian cuisine with its fusion of traditional spices and modern plating. Their signature dish, saffron-infused lamb biryani, has become a favorite among food critics and tourists alike.

Predicted Category: food & drink

Confidence Score: 0.7609241667096004

Final Conclusion & Future Work

In this project, I successfully built a machine learning model to classify news articles into predefined categories using TF-IDF features and classical algorithms. Logistic Regression emerged as the best-performing model, achieving an accuracy of **76%**, while also offering strong interpretability through feature importance analysis.

While the model performs well overall, it struggles slightly with underrepresented categories such as *Business* and *Sports*. This highlights a limitation in class balance and feature richness. Future work could explore the following enhancements:

- **Advanced Models:** Implement deep learning architectures like LSTM or BERT for richer text understanding.
- **Improved Feature Engineering:** Integrate word embeddings such as GloVe or contextual embeddings for semantic depth.
- **Class Imbalance Handling:** Apply techniques like SMOTE or class-weighted loss functions to improve minority class performance.
- **Model Deployment:** Build a user-friendly interface or API for real-time article classification.
- **Explainability Tools:** Use SHAP or LIME to visualize and explain model predictions for transparency.

Overall, this project lays a strong foundation for scalable, automated news categorization. With further enhancements, it can significantly improve content delivery and user experience across digital platforms.

Prepared By: Vandana Bhavsar